

Restoration-Guided State-Space Adaptation for Degradation-Robust RGB Object Tracking

Abdel Rahman Idrais, *Student Member, IEEE*, Alireza Esmaeilzahi, *Member, IEEE*, M. Omair Ahmad, *Fellow, IEEE*, and M. N. S. Swamy, *Fellow, IEEE*

Abstract—Visual object trackers can degrade under image-quality corruptions such as blur, low resolution, and sensor noise. This paper studies a conservative restoration-guided state-space adaptation of OTrack for RGB template-search tracking [1]. The final Paper Freeze V1 method inserts a Restoration-Guided State-Space Block (RG-SSB) into OTrack, keeps the backbone frozen, and trains only `rgssb.*` and `box_head.*` using degradation-aware LaSOT training with global clean-degraded feature consistency at weight $\lambda = 0.02$ [2]. Response consistency, target-region feature consistency, and target-versus-distractor margin loss are not retained in the final method. On selected benchmark subsets, the method improves average AUC on broader OTB by +0.033144 and expanded UAV123 by +0.017405, while corrected expanded NFS is slightly negative at -0.009638. The cross-benchmark aggregate is near-neutral positive at +0.001478. These results support a cautious conclusion: restoration-guided feature adaptation shows robustness potential, but it does not establish state-of-the-art or universal degradation robustness.

Index Terms—Visual object tracking, degradation robustness, state-space models, Mamba, feature consistency, RGB tracking.

I. INTRODUCTION

Single-object tracking is a fundamental component of video analysis systems that must continuously localize a target after it has been specified in an initial frame. In practical deployments, the tracker is expected to operate under camera motion, target deformation, background clutter, partial occlusion, and changes in imaging quality. Benchmarks such as OTB, UAV123, LaSOT, and NFS have therefore become important tools for evaluating localization behavior under diverse sequence conditions [2]–[5]. For RGB template-search tracking, however, the problem is not only to model target appearance over time, but also to preserve discriminative target information when the search image is degraded by blur, reduced resolution, noise, or other acquisition artifacts.

Abdel Rahman Idrais, M. Omair Ahmad, and M. N. S. Swamy are with the Department of Electrical and Computer Engineering, Concordia University, Montreal, QC H3G 1M8, Canada.

Alireza Esmaeilzahi is with The Edward S. Rogers Sr. Department of Electrical and Computer Engineering, University of Toronto, Toronto, ON M5S 3G8, Canada.

Corresponding author: Abdel Rahman Idrais. Email structure requires confirmation; local InvTrack text extraction renders the address as a idrais@encs.concordia.ca.

[AUTHOR INFORMATION REQUIRES CONFIRMATION] Candidate author/funding information is copied from the local InvTrack PDF and requires human confirmation for this manuscript. ORCIDs, biographies, and final submission metadata remain unresolved. This work was supported in part by the Natural Sciences and Engineering Research Council of Canada, and in part by the Regroupement Stratégique en Microélectronique du Québec.

Recent one-stream trackers have shown that joint template-search feature learning can provide an effective and compact tracking formulation. OTrack is a representative example, since it learns relations between the template and search regions within a unified feature stream and is used as the baseline architecture in this work [1]. Such trackers can be effective under clean or moderately challenging observations, but degraded search frames may alter the feature statistics on which template-search matching and box regression depend. This vulnerability is especially relevant when the tracker is trained and evaluated under clean and synthetically degraded branches, because the model must localize the same target even when the degraded search feature is less consistent with its clean counterpart.

Several lines of work are relevant to this robustness problem. Degradation-aware tracking can be formulated through clean/degraded training and consistency constraints, as shown by the local InvTrack evidence, which combines invariant feature learning with response-map fusion [6]. Image restoration methods, including recent state-space restoration models, also aim to recover useful visual structure from degraded observations [7], [8]. These directions are valuable, but they do not directly answer the narrower question considered here. A tracker can be made degradation-aware without using restoration-oriented state-space adaptation, and an image restoration model can improve image quality without being optimized for template-search localization.

State-space models and Mamba-based vision modules have recently attracted attention because they provide a mechanism for structured long-range feature modeling with favorable computational properties [9]–[11]. In visual tracking, reported Mamba-based designs use state-space modules for purposes such as long-term context modeling, multimodal feature interaction, nonlinear motion prediction, dynamic template generation, or nighttime template-search representation learning [12]–[16]. In parallel, MambaIR and MambaIRv2 show that state-space modeling can be adapted to restoration tasks through restoration-oriented feature processing [7], [8]. The unresolved gap addressed in this paper is therefore specific: whether a restoration-guided state-space block can be inserted into an RGB template-search tracker to improve clean-degraded feature consistency, without replacing the tracking framework or claiming a general solution to degradation-robust tracking.

To examine this question under controlled conditions, this paper introduces a Restoration-Guided State-Space Block (RG-SSB) into OTrack. The module is placed after the

frozen backbone search features and before the box head, and the final training configuration updates only `rgssb.*` and `box_head.*`. The method uses global clean-degraded feature consistency with weight $\lambda = 0.02$, while response consistency, target-region feature consistency, target-versus-distractor margin loss, degradation tokens, memory update, response fusion, and template-guided scan are not retained in the final method. This design deliberately separates restoration-guided feature adaptation from image reconstruction, from ordinary preprocessing, and from broader changes to the tracking architecture.

The main contributions of this work are as follows:

- A conservative OSTRack+RG-SSB tracking adaptation is introduced, in which restoration-guided state-space processing is integrated between the frozen backbone features and the box head. This isolates the effect of the proposed block from full-backbone retraining and broad tracker redesign.
- A degradation-aware training objective is evaluated using global clean-degraded feature consistency at $\lambda = 0.02$. The final configuration is supported by the method-selection evidence, while response consistency, target-region feature consistency, and target-versus-distractor margin loss are reported as rejected alternatives rather than components of the proposed method.
- The method is evaluated under clean, motion-blur, low-resolution, and Gaussian-noise conditions on selected OTB, UAV123, and corrected NFS subsets [3]–[5]. The evidence supports positive average AUC changes on the selected OTB and UAV123 subsets, but corrected NFS remains mixed/slightly negative, so the claims are limited to a cautious proof of concept.
- The experimental record includes corrected NFS annotation provenance, same-GPU efficiency and parameter comparisons, and explicit claim boundaries. These checks prevent the corrected results from being combined with superseded NFS metrics and avoid state-of-the-art or universal robustness claims.

The remainder of this paper is organized as follows. Section II reviews related work and positions the proposed method. Section III describes the RG-SSB formulation and training objective. Sections IV and V present the experimental protocol and results, and Sections VI and VII discuss limitations and conclude the paper.

II. RELATED WORK

A. One-Stream and Transformer-Based Visual Tracking

Modern RGB trackers have increasingly used transformer-based feature interaction to relate the initial target template to the search region [17], [18]. One-stream tracking further simplifies this formulation by learning template-search relations in a unified stream rather than maintaining separate feature extraction and matching stages [1], [19]. OSTRack follows this direction and provides the base architecture used in this paper [1]. Its one-stream formulation is well suited to the present study because it gives a fixed template-search tracking

pipeline into which a feature-adaptation block can be inserted without redefining the tracker.

The strength of this family is that target localization can be learned through direct template-search interaction. At the same time, the same dependence on search-region features can make the tracker sensitive to image-quality changes that alter the representation before localization. The present work does not propose a new one-stream tracker, a new temporal-memory mechanism, or a new template-update strategy. Instead, it asks whether the search features produced by an existing one-stream tracker can be made more consistent across clean and degraded observations through a restoration-guided state-space block.

B. Robust Tracking Under Degraded Observations

Robust tracking under degraded observations has been studied from several viewpoints, including low-light benchmark construction, nighttime UAV tracking, degradation-aware training, invariant feature learning, and evaluation under corrupted visual inputs [6], [16], [20]. The local InvTrack evidence is particularly relevant because it studies degradation-invariant single-object tracking with clean/degraded branches, feature consistency, and response-map fusion [6]. It therefore prevents the present manuscript from claiming that degradation-aware tracking, clean-degraded training, or feature consistency is new.

The limitation addressed here is narrower. InvTrack provides direct evidence that degradation-invariant tracking can be formulated without restoration-oriented state-space feature adaptation. Conversely, an adaptation inserted into OSTRack must be evaluated separately because its effect depends on where it is placed, which parameters are trained, and whether the consistency objective improves localization rather than only feature agreement. The present method therefore keeps the backbone frozen, trains RG-SSB and the box head, and reports both positive and negative benchmark behavior instead of treating degradation-aware training as sufficient by itself.

C. State-Space and Mamba-Based Visual Tracking

State-space and Mamba-based modules have been adopted in vision because they can model structured dependencies without relying solely on quadratic attention [9]–[11]. In tracking, recent Mamba-based methods have been reported for roles such as temporal modeling, multimodal fusion, motion reasoning, template update, or efficient feature processing [12]–[16]. These directions are important because they show that state-space modeling is already relevant to tracking, and they also prevent the present work from making a broad novelty claim about using Mamba in tracking.

The present paper is positioned differently. It does not introduce a Mamba tracker with temporal memory, response-map fusion, or a dynamic template-update module. It also does not claim that existing Mamba trackers ignore all forms of degradation, since that would require a complete and separately cited survey. The specific gap retained here is the missing connection between restoration-oriented state-space feature adaptation and RGB template-search tracking under controlled clean/degraded training.

D. Restoration-Oriented State-Space Models

Image restoration methods aim to recover useful image structure under degradations such as blur, noise, and resolution loss. MambaIR adapts state-space modeling to restoration through residual state-space processing, local enhancement, and channel attention [7]. MambaIRv2 further develops attentive state-space restoration and semantic guided neighboring for restoration tasks [8]. These works are technically relevant because they show how state-space modeling can be organized around degraded visual evidence rather than only around recognition or tracking objectives.

However, restoration models and tracking models optimize different outputs. Image restoration typically aims to improve reconstructed or enhanced visual content, whereas template-search tracking must preserve target-discriminative features that support localization. A restoration module used as preprocessing may also change image appearance without guaranteeing better matching between the template and search features. The proposed RG-SSB therefore uses restoration-guided state-space processing inside the tracker feature path, while the training objective remains tied to tracking and global clean-degraded feature consistency.

E. Positioning of the Proposed Method

The proposed method is best understood as a constrained adaptation of OTrack, not as a replacement for existing tracking frameworks. Relative to ordinary one-stream tracking, it adds a search-feature adaptation block and a clean-degraded feature-consistency term. Relative to image restoration preprocessing, it does not reconstruct an enhanced RGB image and does not optimize an image-quality objective. Relative to generic Mamba vision backbones, it uses a restoration-guided state-space block at a specific point in an existing tracker. Relative to degradation-invariant tracking methods such as InvTrack, it avoids response-map fusion and evaluates at a different residual state-space feature-adaptation path [6].

This positioning also defines the negative boundaries of the final method. Response consistency, target-region consistency, target-versus-distractor margin loss, degradation tokens, memory update, response fusion, and template-guided scan were considered outside the retained configuration or rejected by the method-selection evidence. Consequently, the paper makes a limited claim: RG-SSB provides a controlled test of restoration-guided state-space adaptation for OTrack under selected clean and degraded RGB tracking conditions. It does not establish state-of-the-art tracking performance, universal degradation robustness, or uniform improvement across all benchmarks.

III. PROPOSED METHOD

A. Overview

The final method uses OTrack as the base tracker. Given a template image and a search image, OTrack extracts transformer features and predicts a response map and bounding box. Paper Freeze V1 adds a Restoration-Guided State-Space Block (RG-SSB) after the backbone search features and before

the tracking head. The module is designed to adjust tracking features under degradation rather than reconstruct RGB images.

The frozen method is:

- base tracker: OTrack;
- added module: RG-SSB;
- frozen parameters: backbone;
- trainable parameters: `rgssb.*` and `box_head.*`;
- training data: degradation-aware LaSOT [2];
- feature consistency: global clean-degraded feature consistency;
- feature-consistency weight: $\lambda = 0.02$;
- disabled components: response consistency, target-region feature consistency, target-versus-distractor margin loss, degradation tokens, memory update, response fusion, and template-guided scan.

B. Training Objective

The training objective combines the degraded-branch tracking loss with global clean-degraded feature consistency:

$$\mathcal{L} = \mathcal{L}_{\text{track}}(X_d) + 0.02 \mathcal{L}_{\text{feat}}(F_c, F_d), \quad (1)$$

where X_d is the degraded search branch, and F_c and F_d are clean and degraded post-RG-SSB search features. The clean feature branch is detached in the consistency term. The final method intentionally keeps the response-consistency, target-region feature-consistency, and target-versus-distractor margin (TDM) objectives disabled.

IV. EXPERIMENTAL SETUP

A. Baseline and Final Configuration

The baseline is original OTrack config `vitb_256_mae_ce_32x4_ep300`. The final method uses `vitb_256_mae_ce_32x4_ep300_rgssb_head_train_lasot_degraded_hpc_featcons_lam002`. All reported changes are computed as RG-SSB minus the OTrack baseline for the same sequence, degradation, severity, and seed.

B. Benchmarks and Degradation Conditions

The evaluated subsets are:

- broader OTB: 13 sequences and 52 sequence-condition pairs [3];
- expanded UAV123: 16 sequences and 64 pairs [4];
- corrected expanded NFS: 32 sequences and 128 pairs [5].

Each sequence is evaluated under four conditions: clean, motion blur at medium severity, low resolution at medium severity, and Gaussian noise at medium severity. Degraded conditions use seed 42.

C. Metrics

The paper reports Success AUC, Precision@20, mean center error, model parameters, controlled FPS, per-frame latency, and peak allocated GPU memory. FLOPs/MACs are unavailable and are not reported as measured.




Fig. 1. Benchmark-level AUC change for Paper Freeze V1. Positive values favor RG-SSB. The evaluated scope is selected OTB, expanded UAV123, corrected expanded NFS, and the cross-benchmark aggregate.

D. Corrected NFS Provenance

NFS required a correction before current results could be used. Raw NFS annotations were confirmed to contain XXY coordinates, while the previous invalid path treated aligned rows as XYWH. The invalid NFS rows were archived, normalized aligned XYWH annotations were prepared, and NFS tracking was rerun. The current NFS metrics and failure analysis use only corrected normalized aligned XYWH annotations. Earlier NFS metrics and qualitative overlays are superseded.

V. RESULTS

A. Main Benchmark Results

Table I summarizes the benchmark-level results, and Fig. 1 visualizes the same average AUC changes. The strongest result is on OTB, where RG-SSB improves average AUC by +0.033144 across 52 sequence-condition pairs. UAV123 is also positive, with +0.017405 average AUC over 64 pairs. Corrected NFS is the main weakness: the method decreases average AUC by -0.009638 over 128 pairs. The combined cross-benchmark aggregate is +0.001478, which is best described as near-neutral positive rather than consistent superiority.

B. Condition-Level Results

Table II and Fig. 2 summarize condition-level behavior. The method is most consistently positive under motion blur. Low resolution remains the clearest condition-level weakness, with -0.008275 average AUC and increased center error.

C. Ablation Studies

Table III reports the final method-selection ablations, and Fig. 3 visualizes the decisions. Global feature consistency with $\lambda = 0.02$ is retained. Response consistency, target-region feature consistency, and TDM are rejected. TDM is important negative evidence: it improves several local sequence-condition pairs, but its Gaussian-noise regression dominates




Fig. 2. Condition-level AUC change across clean, motion blur, low resolution, and Gaussian noise. The low-resolution condition remains mixed/negative.




Fig. 3. Ablation decision summary. TDM is reported only as rejected negative evidence and is not part of Paper Freeze V1.

the aggregate gate result. It is therefore not part of the frozen method.

The final TDM gate decision is REJECT_AND_ROLL_BACK_TO_PAPER_FREEZE_V1.

D. Efficiency Analysis

Table IV and Fig. 4 report the same-GPU controlled efficiency comparison. The measurement uses a Tesla V100-PCIE-32GB, the same Car1 clean sequence, 50 warm-up frames, 3 repetitions, synchronized CUDA timing, and separate initialization/per-frame tracking measurements. Historical timing files are not mixed with this controlled comparison.

The final method adds 2,079,936 parameters over the baseline. FLOPs/MACs are unavailable and are not reported as measured.

E. Corrected NFS Failure Analysis

Figure 5 summarizes terminal outcomes for 10 corrected NFS failure cases selected from the largest corrected drops. Persistent RG-SSB-specific target loss appears in 5 cases, sudden center jump in 3 cases, shared tracker failure in 1 case,

TABLE I
PAPER FREEZE V1 BENCHMARK SUMMARY. POSITIVE AUC AND PRECISION@20 CHANGES FAVOR RG-SSB; NEGATIVE CENTER-ERROR CHANGES INDICATE LOWER ERROR.

Benchmark	Sequences	Pairs	Baseline AUC	RG-SSB AUC	AUC Change	Precision Change	Center Error Change	Conclusion
OTB	13	52	0.646964	0.680108	+0.033144	+0.050250	-4.563202	Improves clearly on average
UAV123	16	64	0.660183	0.677587	+0.017405	+0.029983	-3.500185	Improves slightly / near neutral
NFS	32	128	0.681149	0.671511	-0.009638	-0.006937	+3.705562	Mixed / negative average
Cross-benchmark aggregate	53	212	0.689269	0.690748	+0.001478	+0.007843	+1.161416	Near-neutral positive overall

TABLE II
CONDITION-LEVEL ROBUSTNESS SUMMARY ACROSS SELECTED OTB, EXPANDED UAV123, AND CORRECTED EXPANDED NFS SUBSETS.

Condition	Pairs	Baseline AUC	RG-SSB AUC	AUC Change	Precision Change	Center Error Change	Conclusion
Clean	61	0.685280	0.690648	+0.005368	+0.015220	+2.558046	Improves on average
Motion blur	61	0.643687	0.663299	+0.019612	+0.029967	-7.502857	Improves on average
Low resolution	61	0.672450	0.664175	-0.008275	-0.003989	+6.769353	Mixed / negative
Gaussian noise	61	0.672042	0.681627	+0.009585	+0.018540	-0.342091	Improves on average

TABLE III
ABLATION SUMMARY FOR PAPER FREEZE V1. LOCAL ABLATIONS GUIDE METHOD SELECTION AND ARE NOT FULL BENCHMARK CLAIMS.

Ablation	Scope	Primary Comparison	Average AUC Change	Decision	Evidence
Global feature consistency $\lambda = 0.02$	3 OTB sequences $\times$ 4 conditions	vs. OTrack baseline	+0.023480	Kept	Lambda sweep CSV
Response consistency	3 OTB sequences $\times$ 4 conditions	vs. feature-only $\lambda = 0.02$	-0.019360	Rejected	Response-consistency CSV
Target-region feature consistency	3 OTB sequences $\times$ 4 conditions	vs. global feature consistency	-0.023555	Rejected	Target-vs-global CSV
Target-vs-distractor margin	3 OTB sequences $\times$ 4 conditions	vs. global feature consistency	-0.014549	Rejected by gate	TDM gate CSV

TABLE IV
CONTROLLED SAME-GPU EFFICIENCY COMPARISON ON TESLA V100-PCIE-32GB. FLOPS/MACs WERE UNAVAILABLE AND ARE NOT REPORTED AS MEASURED.

Model	Params	Trainable Params	Frozen Params	Init. ms	Latency ms	FPS	Peak Alloc. MB	Peak Reserved MB	FLOPs/MACs
OTrack baseline	92,518,533	92,518,533	0	2.949	10.724	93.247	380.233	434.000	Unavailable
RG-SSB final	94,598,469	8,554,053	86,044,416	3.165	11.525	86.766	388.917	434.000	Unavailable




Fig. 4. Controlled V100 runtime comparison. The final method has a modest FPS reduction relative to the baseline and a small increase in peak allocated memory.




Fig. 5. Corrected NFS failure-mode distribution for 10 selected failure cases. Ground truth and diagnostics use corrected aligned XYWH annotations.

and ambiguous behavior in 1 case. This is a qualitative failure-focused analysis, not a benchmark-wide frequency estimate.

VI. DISCUSSION AND LIMITATIONS

The evidence supports a narrow, defensible claim: restoration-guided state-space adaptation can improve selected robustness outcomes in RGB template-search tracking, especially on the tested OTB subset. UAV123 is positive but mixed,

and corrected NFS is slightly negative. The result is therefore promising but not uniformly improved.

The low-resolution weakness and corrected NFS failures are important limitations. They suggest that global feature consistency may not fully protect target identity under fast

motion, high-frame-rate sampling behavior, or severe target/background ambiguity. The rejected TDM ablation further shows that adding a response-margin objective is not automatically beneficial; poorly balanced response constraints can introduce large condition-specific regressions.

This paper version has several boundaries. The benchmark coverage uses selected subsets, not complete benchmark suites. Degradations are synthetic, medium severity, and use one seed. Training uses LaSOT only. The backbone remains frozen. FLOPs/MACs are unavailable. Corrected NFS required annotation normalization and rerunning NFS tracking; older NFS metrics and overlays are superseded. The evidence does not support state-of-the-art claims, universal degradation robustness, or uniform gains across all benchmarks.

VII. CONCLUSION

Paper Freeze V1 establishes a conservative fallback method: OSTRack with RG-SSB and box-head training under global clean-degraded feature consistency. The method is positive on OTB and UAV123, mixed on corrected NFS, and near-neutral positive in the mandatory cross-benchmark aggregate. This is sufficient for a cautious proof-of-concept paper position, but not for claims of broad superiority. Future work should expand benchmark coverage, test multiple seeds and severities, and investigate failure-specific refinements only after the current evidence is presented with appropriate limits.

APPENDIX A REPRODUCIBILITY NOTES

The final method config is `vitb_256_mae_ce_32x4_ep300_rgssb_head_train_lasot_degraded_hpc_featcons_lam002`. The frozen method keeps the backbone frozen and trains only `rgssb.*` and `box_head.*`. NFS results must use the corrected normalized aligned XYWH annotation root. The invalid historical NFS rows are archived for provenance only and must not be used for current paper claims. FLOPs/MACs are unavailable and should not be reported as measured.

REFERENCES

- [1] B. Ye, H. Chang, B. Ma, S. Shan, and X. Chen, "Joint feature learning and relation modeling for tracking: A one-stream framework," in *Computer Vision – ECCV 2022*, ser. Lecture Notes in Computer Science, 2022, pp. 341–357.
- [2] H. Fan, L. Lin, F. Yang, P. Chu, G. Deng, S. Yu, H. Bai, Y. Xu, C. Liao, and H. Ling, "Lasot: A high-quality benchmark for large-scale single object tracking," in *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 5369–5378.
- [3] Y. Wu, J. Lim, and M.-H. Yang, "Online object tracking: A benchmark," in *2013 IEEE Conference on Computer Vision and Pattern Recognition*, 2013, pp. 2411–2418.
- [4] M. Mueller, N. Smith, and B. Ghanem, "A benchmark and simulator for uav tracking," in *Computer Vision – ECCV 2016*, ser. Lecture Notes in Computer Science, 2016, pp. 445–461.
- [5] H. K. Galoogahi, A. Fagg, C. Huang, D. Ramanan, and S. Lucey, "Need for speed: A benchmark for higher frame rate object tracking," in *2017 IEEE International Conference on Computer Vision (ICCV)*, 2017, pp. 1134–1143.
- [6] A. R. Idrais, A. Esmailzahi, M. O. Ahmad, and M. N. S. Swamy, "Invtrack: Efficient deep object tracking neural network with invariant feature learning," local InvTrack PDF verifies title, authors, affiliations, corresponding author, and funding text, but contains placeholder publication dates and DOI. Use only as local evidence until final publication metadata is confirmed.

- [7] H. Guo, J. Li, T. Dai, Z. Ouyang, X. Ren, and S.-T. Xia, "Mambair: A simple baseline for image restoration with state-space model," in *European Conference on Computer Vision*, 2024.
- [8] H. Guo, Y. Guo, Y. Zha, Y. Zhang, W. Li, T. Dai, S.-T. Xia, and Y. Li, "Mambairv2: Attentive state space restoration," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025.
- [9] A. Gu and T. Dao, "Mamba: Linear-time sequence modeling with selective state spaces," in *First Conference on Language Modeling*, 2024.
- [10] L. Zhu, B. Liao, Q. Zhang, X. Wang, W. Liu, and X. Wang, "Vision mamba: Efficient visual representation learning with bidirectional state space model," in *Proceedings of the 41st International Conference on Machine Learning*, 2024.
- [11] Y. Liu, Y. Tian, Y. Zhao, H. Yu, L. Xie, Y. Wang, Q. Ye, J. Jiao, and Y. Liu, "Vmamba: Visual state space model," 2024.
- [12] X. Li, B. Zhong, Q. Liang, G. Li, Z. Mo, and S. Song, "Mambalot: Boosting tracking via long-term context state space model," 2024.
- [13] J. Huang, S. Wang, S. Wang, Z. Wu, X. Wang, and B. Jiang, "Mambaftrack: Frame-event tracking via state space model," 2024.
- [14] C. Xiao, Q. Cao, Z. Luo, and L. Lan, "Mambatrack: A simple baseline for multiple object tracking with state space model," in *Proceedings of the 32nd ACM International Conference on Multimedia*, 2024.
- [15] X. Wang, C. Wang, S. Wang, X. Wang, Z. Zhao, L. Zhu, and B. Jiang, "Mambaevt: Event stream based visual object tracking using state space model," 2024.
- [16] Y. Wu, X. Yang, X. Wang, H. Ye, D. Zeng, and S. Li, "Mambanut: Nighttime uav tracking via mamba and adaptive curriculum learning," 2024.
- [17] X. Chen, B. Yan, J. Zhu, D. Wang, X. Yang, and H. Lu, "Transformer tracking," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021.
- [18] B. Yan, H. Peng, J. Fu, D. Wang, and H. Lu, "Learning spatio-temporal transformer for visual tracking," in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021.
- [19] Y. Cui, C. Jiang, L. Wang, and G. Wu, "Mixformer: End-to-end tracking with iterative mixed attention," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022.
- [20] P. Zhong, X. Guo, D. Huang, X. Peng, Y. Li, Q. Zhao, and S. Li, "Low-light object tracking: A benchmark," 2024.